

Lab 3 Report: Spark-Based Data Lake

Group L3-T04

Marvin Ernst, Oriol Gelabert, Alex Malo

June 24, 2025

Overview of Pipeline Architecture

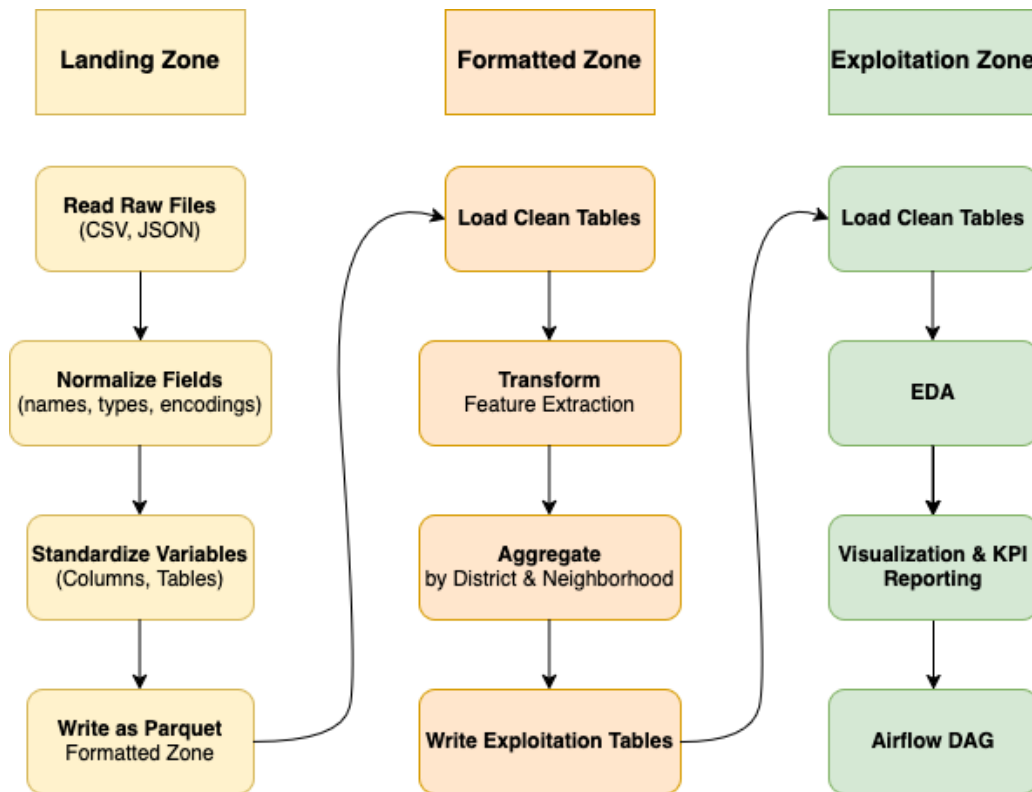


Figure 1: High-level abstraction of the three main pipelines

- **Pipeline 1: Data Formatting** (Raw → Cleaned Parquet)

- Input: CSV/JSON files from `landing_zone/`
- Tasks: unzip, harmonize schema, type casting, renaming, spatial normalization
- Output: Cleaned and harmonized Parquet files in `formatted_zone/`

- **Pipeline 2: Exploitation** (Cleaned → Aggregated)

- Input: Parquet files from `formatted_zone/`
- Tasks: aggregations at multiple spatial levels (district, neighborhood, census section), KPI computation
- Output: Indicator datasets in `exploitation_zone/`

- **Pipeline 3: Analysis** (Aggregated → Dashboards)
 - Input: Aggregated indicator files
 - Tasks: descriptive analysis and plots using Seaborn and Matplotlib
 - Output: Visual dashboards to reveal trends and spatial patterns

Selected Datasets and Assumptions

We selected the following three datasets, all located in `landing_zone/`:

1. **Tourist Housing**: quarterly CSV files with apartment registration data
2. **Commercial Premises**: yearly CSV files on ground-floor businesses
3. **Household Size**: JSON files with household structure (number of people per dwelling)

Key assumptions made:

- **District/Neighborhood Mapping**: Many datasets lacked district and neighborhood codes but included names. We extracted reliable mappings from the 2023 tourist housing files and built mapping dictionaries.
- **Harmonization**: Yearly and quarterly datasets had inconsistent schemas. We renamed and typecast all relevant columns for cross-year integration.
- **Boolean Fields**: Strings like "Sí"/"No" or other variations were normalized to binary indicators for commercial premises.
- **Date Inference**: Registration dates were inferred from string IDs (e.g., `N_EXPEDIENT` in housing).
- **Missing Data**: Null values were preserved unless necessary to replace (e.g., for joins or aggregations).

Pipeline Specific Design Justifications

Data Formatting Pipeline

- Used regex and encoding detection to process CSV and JSON files robustly
- Applied batch extraction, delimiter correction, and header normalization to ensure Spark compatibility
- Used UDFs to normalize accent marks and string cases for consistent joins
- Final outputs stored in columnar Parquet format for downstream efficiency

Exploitation Pipeline

- Aggregated indicators by district, neighborhood, and section to support flexible spatial analysis
- Computed KPIs like share of coworking/nightlife premises, average household size, and total tourist licenses
- Calculated relative proportions (% of premises or households) instead of raw counts to allow comparability
- Stored all outputs in `exploitation_zone/` with clear file naming convention for traceability

Analysis Pipeline

- Combined Spark (ETL) with Pandas for final summarization and plotting (performance trade-off)
- Produced plots for:
 - Household size distributions
 - Commercial indicators by district
 - Tourist housing trends over time
 - Correlation between tourist activity and commercial premises
- Constructed a KPI summary table merging all indicators

Airflow Orchestration

To ensure reproducibility and automation, we implemented a DAG in Apache Airflow that sequentially runs:

1. `01_data_formatting_pipeline.py`
2. `02_exploitation_pipeline.py`
3. `03_analysis_pipeline.py`

Each script is executable and reads paths from environment variables to avoid hardcoding. Airflow provides a centralized orchestration mechanism ensuring repeatable workflows.